

Variable-Length Message-Passing Graph Neural Networks as Surrogates for 1D Finite Element Bead Chain Simulations

Grey L. Goodwin

*Dept. of Electrical and Computer Engineering
University of Kentucky
Lexington, KY, USA
grey.goodwin@uky.edu*

Daniel L. Lau, PhD.

*Dept. of Electrical and Computer Engineering
University of Kentucky
Lexington, KY, USA
dllau@uky.edu*

Abstract—Finite element analysis (FEA) of a flexible bead chain requires thousands of small, localized force and integration computations per simulated second. We investigate whether a graph neural network (GNN) trained on volume-sampled single-step physics can act as a fast surrogate for the per-node update in a 1D chain of tapered mass nodes connected by spring rod elements. An initial shared multilayer perceptron trained on a fixed 16-bead chain achieved accurate single-step predictions but could not generalize to chains of other lengths. We address this by restructuring the network as a true message-passing GNN with separate encoders for node and edge features, ghost masking for chain endpoints, and weight sharing across all nodes. The same trained model runs on chains of any length. We train on 1M per-bead samples drawn from chains of 8 to 32 beads with random taper ratios, and evaluate rollout accuracy, constraint satisfaction, and generalization on chain lengths never seen during training. The model maintains accurate behavior across the trained range and produces qualitatively correct dynamics under varied initial orientations and velocities. We discuss the connection between the local message-passing structure and the underlying physics, the relationship between sampling strategy and generalization, and remaining limitations around long-horizon drift and exact constraint satisfaction.

Index Terms—graph neural network, finite element analysis, message passing, surrogate model, bead chain, volume sampling

I. INTRODUCTION

Physical simulations of flexible one-dimensional structures such as whips, ropes, cables, and biological filaments are widely used in mechanical engineering, aerospace, and animation. The canonical formulation discretizes the structure as a chain of mass nodes connected by rod elements, applies local force laws, and integrates forward in time. Even simple cases involve small timesteps and many integration steps, making long simulations expensive despite the low dimensionality of each timestep.

Machine learning surrogates for physics simulations have attracted attention as a potential source of speedup. Graph neural networks (GNNs) are a natural fit for mesh-based simulations because the mesh itself is a graph: nodes are discretization points, edges are elements, and the local physics at each node depends only on its immediate neighbors. A

trained GNN that respects this graph structure can, in principle, replace the expensive per-timestep numerical computation with a single forward pass.

The project we describe here explores this question on a 1D FEA bead chain that models a whip. The chain has N mass nodes connected by $N - 1$ spring rod elements, with a tapered mass distribution from a heavy fixed anchor to a light free tip. Each bead’s next state depends only on its own state and the states of its immediate neighbors, a property we refer to as locality. The reference solver is a symplectic Euler integrator with optional SHAKE constraint projection for rod inextensibility.

Our initial surrogate, a shared multilayer perceptron (MLP) that takes a concatenated vector of self and neighbor features, achieved sub-percent error on single-step predictions for a fixed 16-bead chain. However, the architecture is tightly coupled to the chain length: the input vector’s dimensionality, the mass distribution seen during training, and the rest length are all fixed by $N = 16$. The model cannot run on chains of other sizes, and any new chain length requires regenerating data and retraining from scratch.

This paper describes the architectural change we made to remove that limitation. We restructure the surrogate as a true message-passing GNN with separate node and edge encoders, message functions applied per edge, and ghost masking for chain endpoints. Because all learned functions have shared weights applied per-node, the same trained model handles chains of any length. We train on per-bead samples drawn from chains of varying lengths and taper ratios, and we verify that the model generalizes to chain lengths not seen during training.

The contributions of this paper are:

- 1) A variable-length message-passing GNN architecture for 1D bead chain physics, compatible with per-bead volume-sampled training data.
- 2) A data generation procedure that samples single-step physics scenarios uniformly across chain lengths, taper ratios, and local states, enabling training without trajectory rollouts.

- 3) An empirical evaluation of the resulting surrogate on chain lengths, initial orientations, and initial velocities that were not seen during training.

II. BACKGROUND AND PROBLEM FORMULATION

A. Bead Chain Physics

The chain consists of N mass nodes connected by $N - 1$ rod elements. Node 0 is fixed at an anchor point; nodes $1, \dots, N - 1$ are free. Each bead has mass m_i , position $\mathbf{x}_i \in \mathbb{R}^2$, and velocity $\mathbf{v}_i \in \mathbb{R}^2$. Mass is distributed with a linear taper controlled by a ratio $r = m_0/m_{N-1}$; $r = 10$ produces a whip-like gradient from heavy anchor to light tip.

Three forces act on each free bead: gravity, rod spring forces from each connected neighbor (Hooke's law), and optional viscous damping along the rod axis. A small global velocity drag provides gentle settling.

For rod element e connecting beads i and j with rest length ℓ_0 , the spring force magnitude is

$$\mathbf{F}_{ij}^{\text{spring}} = k(d - \ell_0)\hat{\delta}, \quad (1)$$

where $d = |\mathbf{x}_j - \mathbf{x}_i|$, $\hat{\delta} = (\mathbf{x}_j - \mathbf{x}_i)/d$, and k is the spring stiffness. A viscous damping force $\mathbf{F}_{ij}^{\text{damp}} = c((\mathbf{v}_j - \mathbf{v}_i) \cdot \hat{\delta})\hat{\delta}$ resists relative motion along the rod axis. Total force on bead i is the sum of gravity, drag, and contributions from left and right rods.

Time integration uses symplectic (semi-implicit) Euler:

$$\mathbf{v}_i^{n+1} = \mathbf{v}_i^n + \Delta t \mathbf{a}(\mathbf{x}^n) \quad (2)$$

$$\mathbf{x}_i^{n+1} = \mathbf{x}_i^n + \Delta t \mathbf{v}_i^{n+1} \quad (3)$$

The key property is that the position update uses the *new* velocity, not the old one, which preserves the symplectic form of Hamiltonian mechanics and prevents long-term energy drift.

For the experiments in this paper we use $k = 10^4$ N/m, $c = 0.5$, $\Delta t = 10^{-4}$ s, $N = 16$ unless otherwise stated, total chain length $L = 2.0$ m, total mass $M = 0.5$ kg, and taper ratio $r = 10$. These are default values for the reference solver.

B. Graph Structure of the Problem

The finite element mesh is naturally a graph. Nodes are beads with features $(\mathbf{v}, m, \text{is_fixed})$; edges are rods with features $(\Delta \mathbf{x}, \Delta \mathbf{v}, m_{\text{nbr}}, \ell_0)$ where the Δ quantities are neighbor-relative. Each interior bead has exactly two neighbors (left and right). Endpoint beads (bead 0 and bead $N - 1$) have only one neighbor each.

For one timestep, the physical interaction is strictly 1-hop: the force on bead i depends only on $\mathbf{x}_i$, $\mathbf{v}_i$, m_i , and the equivalent quantities for its immediate neighbors $i - 1$ and $i + 1$. No 3-body or longer-range interaction exists. This locality is central to the architectural choices described later.

C. Initial MLP Surrogate and Its Limitation

Our first surrogate was a shared MLP that took a 16-dimensional concatenated feature vector per bead: 4 self features $[v_x, v_y, m, \text{is_fixed}]$, 6 left-neighbor features $[\Delta x, \Delta y, \Delta v_x, \Delta v_y, m_{\text{nbr}}, \ell_0]$, and 6 right-neighbor features

with the same structure. The output was the 4-dimensional residual state delta $[\Delta x, \Delta y, \Delta v_x, \Delta v_y]$ to be added to the bead's current state.

This architecture is applied per-bead with shared weights, so it is sometimes informally called a GNN. However, the model has no separation between node and edge features, no explicit message computation, and no mechanism for handling a varying number of neighbors. The input vector's composition is hard-coded by the chain length: the per-bead mass distribution seen during training is specific to a fixed N , and the model cannot accept inputs from a chain of any other size.

For a fixed 16-bead chain, this approach works well in single-step prediction. The cost is that the model is locked to one chain topology. To simulate an 8-bead or 24-bead chain, the data generator and the model must both be rebuilt from scratch. The following section describes the architectural change we made to remove this limitation.

III. METHOD

A. Message-Passing GNN Architecture

We replace the flat MLP with a message-passing GNN (hereafter BeadMPGNN). The architecture consists of five learned components:

- **Node encoder**: a residual MLP that maps the 4-dimensional node feature vector to a 64-dimensional hidden state.
- **Edge encoder**: a separate residual MLP that maps the 6-dimensional edge feature vector (neighbor-relative) to a 64-dimensional hidden state. Used for both left and right edges.
- **Message MLP**: maps an encoded edge to a 64-dimensional message.
- **Node update MLP**: maps the concatenation $[\mathbf{h}_{\text{self}} \parallel \mathbf{m}_{\text{left}} \parallel \mathbf{m}_{\text{right}}] \in \mathbb{R}^{192}$ back to a 64-dimensional updated node state.
- **Output head**: a linear layer that projects the final hidden state to the 4-dimensional residual state delta.

Each bead's forward pass proceeds as follows:

- 1) Normalize input features using stored per-feature mean and standard deviation (computed once from the training set and stored as frozen buffers).
- 2) Encode self state and both edges through their respective encoders.
- 3) Compute left and right messages by passing each edge encoding through the message MLP.
- 4) Zero out messages from absent neighbors via ghost masking: $\mathbf{m}_L \leftarrow \mathbb{I}_{\text{has_left}} \cdot \mathbf{m}_L$, and similarly for right.
- 5) Concatenate $[\mathbf{h}_{\text{self}} \parallel \mathbf{m}_L \parallel \mathbf{m}_R]$ and run through the node update MLP.
- 6) Apply the output head and denormalize to obtain the residual deltas $[\Delta x, \Delta y, \Delta v_x, \Delta v_y]$.
- 7) Add the deltas to the current state to produce the next state.

The total parameter count is 80,324. All weights are shared across all beads in a chain, regardless of chain length.

Two architectural choices enable variable-length operation. First, ghost masking structurally enforces the constraint that absent neighbors contribute nothing. When bead 0 processes its left edge features (which are zero because no left neighbor exists), the edge encoder and message MLP produce some nonzero output; this output is multiplied by the zero mask, forcing it to zero before aggregation. The model does not need to learn that all-zero edge features mean “no neighbor.” Second, all learned functions are weight-shared and applied per-bead. A 32-bead chain uses the same 80K parameters as an 8-bead chain; only the number of per-bead forward passes changes.

For this paper we use one round of message passing ($K = 1$), which matches the 1-hop receptive field of the underlying physics. Deeper architectures are possible but would require multi-hop training data or end-to-end rollout training to be effective.

B. Volume-Sampled Per-Bead Training Data

Rather than generating training data by recording simulation trajectories, we sample single-step physics scenarios uniformly across the reachable state space. Each training sample represents one bead at one instant plus the analytical answer for its state one timestep later. No simulation is ever run during data generation; each sample is a completely independent random scenario solved by direct evaluation of the physics equations.

The procedure for one sample is:

- 1) Draw a chain length $N \sim \text{Uniform}(8, 32)$ and a taper ratio $r \sim \text{Uniform}(1, 10)$.
- 2) Compute the tapered mass distribution for that chain and pick a random bead index $i \in \{0, \dots, N - 1\}$.
- 3) Draw self velocity (v_x, v_y) uniformly from empirically observed trajectory bounds.
- 4) Draw each neighbor’s relative position in polar coordinates: radius near the rest length ℓ_0 (within $\pm 5\%$) and angle uniformly in $[0, 2\pi)$. Draw relative velocities uniformly from observed bounds.
- 5) If the bead has no left or no right neighbor, the corresponding features are set to zero.
- 6) Compute the exact single-step physics (springs, damping, gravity, drag, symplectic Euler) to produce the target state delta.

The output is a dataset of 1,000,000 samples, each with four arrays: node features (4-dim), left-edge features (6-dim), right-edge features (6-dim), binary flags for neighbor presence, and 4-dim target deltas. Fig. 1 summarizes the pipeline visually.

This sampling strategy covers the reachable state space more uniformly than trajectory recording. A trajectory-based dataset consists of correlated states along a thin manifold; during rollout, small prediction errors push the simulation slightly off this manifold, where the model has no training coverage and predictions degrade. Volume sampling broadens the coverage so the model sees a dense sample of the states it can plausibly encounter, regardless of the specific trajectory.

C. Training

We normalize node features and edge features separately, each with its own mean and standard deviation vector computed over the full training set. Targets are normalized the same way. The model trains with mean squared error loss on the normalized deltas.

We use AdamW with weight decay 10^{-5} and a cosine annealing learning rate schedule from 10^{-3} . Batch size is 512 and we train for 200 epochs. A 10% split is reserved for validation. Gradient clipping at norm 1.0 is applied throughout. The normalization statistics are stored as buffers on the trained model so that inference is fully self-contained.

IV. RESULTS

A. Training Convergence

Training loss decreases smoothly over 200 epochs and tracks the validation loss closely, indicating the model is neither overfitting nor underfitting on the 1M sample dataset. A scatter of predicted versus true velocity deltas on 200 held-out samples falls tightly along the diagonal, showing that per-sample predictions are accurate. Fig. 2 shows both plots.

B. Generalization to Unseen Chain Lengths

The most important test is whether the model trained on chains of 8–32 beads can produce accurate rollouts on chain lengths not seen during training. We compare the BeadMPGNN rollout with the reference symplectic Euler solver, starting from identical initial conditions, and measure mean position error across all beads over time.

Fig. 3 shows results for chain lengths $N \in \{8, 10, 12, 16, 20, 24, 32\}$. Short to medium chains ($N \leq 16$) track the reference closely with low accumulated error. Longer chains ($N \geq 20$) show more drift, consistent with the longer chain having more compounding per-step errors and more numerous light-mass tip beads. Constraint satisfaction (maximum relative rod length error) is in the 10^{-3} range across all tested lengths, which is acceptable for the stiff-spring formulation of the reference solver but worse than a full SHAKE projection.

C. Initial-Condition Robustness

We additionally tested the model under initial conditions outside the typical “released from horizontal” default: starting the chain rotated to 45° or 90° from horizontal, and applying uniform initial velocities (a downward kick of $v_y = -5$ m/s, or a sideways throw of $v_x = 8$ m/s). Because the training data was volume-sampled rather than trajectory-sampled, the model was exposed to many states that correspond to such initial conditions even though no specific trajectory from these configurations was recorded. The model produces qualitatively correct dynamics under all tested initial conditions within its training velocity bounds. Pushing velocities significantly outside the sampling bounds (e.g. $v_x > 20$ m/s) produces visible divergence, as expected.

generate_volume_data.py — Pipeline Visualization




Fig. 1. Volume sampling pipeline for generating per-bead training data. (1) The physical system is a tapered chain of 16 beads connected by spring rod elements. (2) Each sample is a random scenario: one bead with left and right neighbors sampled in polar coordinates around rest length. (3) Forces on the central bead are computed analytically. (4) Symplectic Euler advances velocity then position. (5) Input and target arrays are packed. (6) Volume sampling fills the reachable state space uniformly, unlike trajectory data which traces thin paths.




Fig. 2. BeadMPGNN training convergence. Left: train and validation MSE loss over 200 epochs on a log scale. Right: predicted versus true velocity delta Δv_x on 200 validation samples.




Fig. 3. Generalization across chain lengths. Left: mean position error between MPGNN and reference over 20,000 timesteps, by chain length. Center: final position error at $t = 2$ s for each chain length. Right: maximum rod length violation. The model is trained on chains of 8–32 beads with random taper ratios, but specific chain lengths such as $N = 14$ or $N = 18$ are never seen as single-sample configurations during training.

V. DISCUSSION

D. Parameter Efficiency

The MPGNN has 80,324 trainable parameters. By comparison, a naive approach that gave each bead its own dedicated MLP would scale linearly with chain length and could not generalize to new sizes. Weight sharing reduces parameter count and enables size-agnostic inference in the same architectural change.

A. Why Message Passing Matters Here

A prior result within our research group showed that for 1D path graphs with fixed topology, the simple concatenation of left and right neighbor states into one MLP input is sufficient to learn local flux functions in compressible fluid problems. The conclusion was that explicit message passing is unnecessary when the graph is a fixed-size path. That conclusion remains valid in its original setting.

The present work does not contradict that result; it addresses a different requirement. The concatenation approach has a

hard architectural ceiling: its input dimensionality is fixed by the chain size. When the goal is a single model that handles multiple chain lengths, the architecture must respect the graph structure explicitly, with separate node and edge processing and shared weight functions applied per-node. Message passing is the mechanism that accomplishes this. The gain is not accuracy per sample; it is the ability to change N at inference time without retraining.

B. Limitations

a) Rollout drift: The model is trained on single-step targets, not rollout trajectories. At inference, its own predictions become the next step’s input, and any single-step error that is correlated with the direction of motion will compound. We observe this especially on long chains where the light tip bead accumulates the most error. Rollout training, which feeds the model’s own output back during training and penalizes multi-step error, is a natural extension that would reduce this drift but was not pursued in the current work.

b) Constraint satisfaction: The reference solver uses stiff springs as a soft constraint; we train the GNN to match those spring-based dynamics. The resulting rod length errors are in the 10^{-3} range rather than machine-zero, which is adequate for visually correct dynamics but not for applications where exact inextensibility matters. Two approaches could improve this: adding a constraint violation term to the training loss, or running a small number of SHAKE iterations after each GNN step as a correction.

c) Receptive field: We use $K = 1$ rounds of message passing, which matches the locality of one physics timestep. For longer rollouts or more complex graphs, multiple message-passing rounds may be needed. However, effectively training $K > 1$ would require training data that captures multi-hop dependencies, either through subgraph sampling or end-to-end trajectory rollouts. The present volume sampling procedure is strictly 1-hop.

d) Domain specificity: The model is trained on a specific physical setup (bead chain with gravity, spring rods, damping, drag). Transferring it to a different physical system would require regenerating data and retraining. However, the architecture itself is not problem-specific: similar patterns apply to any 1D mesh with local interactions.

VI. RELATED WORK

Graph neural networks for physical simulation have a rich literature. Sanchez-Gonzalez et al. [1] demonstrated GNN surrogates for mesh-based particle systems, including fluids and deformable solids. Pfaff et al. [2] extended this to dynamic meshes. Battaglia et al. [3] introduced the interaction network formulation that underlies many modern physics GNNs, including the architectural choices we use here (separate node and edge encoders, message passing, structured aggregation).

Hamiltonian and Lagrangian neural networks (Greydanus et al. [4]) take a different approach by learning a scalar energy function and deriving dynamics through autograd, which guarantees energy conservation structurally. A parallel research

thread in our group is exploring Hamiltonian GNN variants for this same bead chain problem; that work is complementary to the present paper and targets a different failure mode (long-horizon energy drift) with a different architectural mechanism.

Previous work within our research group on the MeltFlow compressible Euler problem demonstrated that simple concatenation works for 1D fixed-topology meshes, informing the architectural comparison in Section IV-A.

VII. CONCLUSION AND FUTURE WORK

We have described a variable-length message-passing graph neural network that acts as a per-timestep surrogate for a 1D finite element bead chain. The architectural change from a shared MLP to a true GNN was motivated not by in-distribution accuracy but by the need for one model to handle chains of varying length. Combined with volume-sampled training data that broadly covers the reachable state space, the resulting model generalizes across chain lengths, initial orientations, and initial velocities within its training bounds.

Several directions extend this work. First, rollout training and multi-step loss would reduce long-horizon drift by exposing the model to the effects of its own errors during training. Second, embedding exact constraint projection into the inference loop, either through a hybrid GNN plus SHAKE step or through a constraint-aware loss term, would improve rod length fidelity. Third, the current 1-hop model assumes strict locality; extending to $K > 1$ rounds would require multi-hop training data but could improve accuracy on wave propagation phenomena that span multiple beads per timestep. Finally, integrating the MPGNN surrogate into a larger engineering workflow, for example real-time visualization or optimization loops that require many forward simulations, is the target application that motivates the speed-accuracy tradeoff.

The broader lesson from this work is that for mesh-based FEA problems, the right architectural choice depends on what the model must generalize over. If the mesh is fixed, simpler architectures may suffice. If the mesh topology varies at inference time, explicit graph structure in the network is what makes variable-length generalization possible at all.

ACKNOWLEDGMENT

The authors thank the University of Kentucky Department of Electrical and Computer Engineering for support.

REFERENCES

- [1] A. Sanchez-Gonzalez, J. Godwin, T. Pfaff, R. Ying, J. Leskovec, and P. W. Battaglia, “Learning to simulate complex physics with graph networks,” in *Proc. ICML*, 2020.
- [2] T. Pfaff, M. Fortunato, A. Sanchez-Gonzalez, and P. W. Battaglia, “Learning mesh-based simulation with graph networks,” in *Proc. ICLR*, 2021.
- [3] P. W. Battaglia, R. Pascanu, M. Lai, D. J. Rezende, and K. Kavukcuoglu, “Interaction networks for learning about objects, relations and physics,” in *Proc. NeurIPS*, 2016.
- [4] S. Greydanus, M. Dzamba, and J. Yosinski, “Hamiltonian neural networks,” in *Proc. NeurIPS*, 2019.